

Fraser Greenlee-Scott

Email: fraser.greenlee@mac.com Mobile: +447989190626 LinkedIn: linkedin.com/in/fraser-greenlee-9658a7101/
Blog: frasgreen.com GitHub: github.com/Fraser-Greenlee

Work

Principal Software Engineer - Salesforce, London, England (July 2025-Present)

Building RL infrastructure for training GUI agents in simulated desktop environments.

- Developed a Ray-based RL training library scaling to 1k+ parallel environments with 126 steps/hour at batch size 1024
- Created a task generator and evaluator yielding +10% absolute success rate improvements on OS-World
- Led a team of 3 to train a grounding model from non-restricted open-source weights, removing dependency on externally-restricted models

Founding Member of Technical Staff (Research Scientist) - Convergence, London, England (July 2024-July 2025)

Built the foundational codebase for Proxy, the company's flagship web agent product.

- Designed the browser tool specification and Reflexion-style prompting framework that maximised agent performance
- Created our grounding setup enabling VLMs to interact with web elements before coordinate-based grounding was reliable
- Built the evaluation infrastructure for web agents; detailed trajectory analysis drove iterative improvements to our agents
- Developed the training pipeline for GUI agent models, producing 10k+ synthetic trajectories and processing customer trajectories into an SFT dataset with multi-layer automatic quality evaluation at both trajectory and step level
- Trained Proxy-Lite, a 3B VLM achieving SOTA on WebVoyager at its scale (938 GitHub stars); also open-sourced a lightweight Playwright-based web agent framework
- Halved inference costs by compressing the two-step think-then-act loop into a single reasoning step
- Implemented context management enabling agents to solve long-horizon 100 step tasks
- Led setup of distributed vLLM hosting serving 125k+ users; step-level evaluation heuristics were integrated into production to monitor model performance

Machine Learning Engineer - Cohere, London, England (May 2022-July 2024)

Started the Command project, a generic instruction following LLM which is now Cohere's primary model.

- Produced core instruction following and preference datasets with external and internal annotation teams.
- Led human preference evaluations to benchmark our models against competitors. Ran multiple evaluations weekly, with results presented to the board to demonstrate competitive parity.
- Created an automated human evaluation tool using reward models, which became the company standard with hundreds of auto-evals conducted across teams.
- Created training infrastructure using JAX used to create state-of-the-art reward and vision models.
- Implemented a shared tokenizer used across Cohere's internal repos and resolved issues with the production tokenizer.
- Research - Spearheaded improvements in reward models, achieving SOTA results on several benchmarks including RewardBench where we outperformed GPT-4 and other top LLMs in preference modelling. Co-authored the paper "Improving Reward Models with Synthetic Critiques".

Machine Learning Engineer - Darktrace, Cambridge, England (March 2021-May 2022)

As part of the Developer-Analyst team, I made ML models & services to alert customers on indicators of compromise.

- Led replacement of old heuristic-based classifiers to using transformers, categorising hostnames, file paths, emails & more.

Software Engineer - Skyscanner, Edinburgh, Scotland (Sept 2019-March 2021)

Education

University of Glasgow - BSc Software Engineering, 1st Class (2015-2019)

Projects/Awards

Hugging Face JAX Contest (June 2021)

Top 15 in Hugging Face Flax-JAX Community Week (2021) trained a T5-based MMD-VAE on TPUs

Datasets of state changes. (2019-2020)

Developed a dataset of Python state changes, crucial for the groundbreaking "[Show Your Work: Scratchpads for Intermediate Computation with Language Models](#)" paper by Google Research.

Pioneer Award - Winner (2019)

Created a [VSCode Extension](#) which showed how your code executes as you type getting over 27k installs.